

ECE 57000 Assignment 07 Exercise

Your Name: Mohmmad Alwakeel

For this assignment, you will do an ablation study on the DCGAN model discussed in class and implemented WGAN with weight clipping and (optional) WGAN with gradient penalty.

Exercise 1: Ablation Study on DCGAN

An [ablation study](https://en.wikipedia.org/wiki/Ablation_(artificial_intelligence)) ([https://en.wikipedia.org/wiki/Ablation_\(artificial_intelligence\)](https://en.wikipedia.org/wiki/Ablation_(artificial_intelligence))) measures performance changes after changing certain components in the AI system. The goal is to understand the contribution on each component for the overall system.

Task 1.0 Original DCGAN on MNIST from class note

Here is the copy of the code implementation from [course website](https://www.davidinouye.com/course/ece57000-fall-2021/lectures/dcgan-mnist-edit.pdf) (<https://www.davidinouye.com/course/ece57000-fall-2021/lectures/dcgan-mnist-edit.pdf>). Please run the code to obtain the result and **use it as a baseline to compare the results** with the following the ablation tasks.

Hyper-parameter and Dataloader setup

```
In [2]: from __future__ import print_function
        %matplotlib inline
        import argparse
        import os
        import random
        import torch
        import torch.nn as nn
        import torch.nn.parallel
        import torch.backends.cudnn as cudnn
        import torch.optim as optim
        import torch.utils.data
        import torchvision.datasets as dset
        import torchvision.transforms as transforms
        import torchvision.utils as vutils
        import numpy as np
        import matplotlib.pyplot as plt
        import matplotlib.animation as animation
        from IPython.display import HTML

        # Set random seed for reproducibility
        manualSeed = 999
        #manualSeed = random.randint(1, 10000) # use if you want new results
        print("Random Seed: ", manualSeed)
        random.seed(manualSeed)
        torch.manual_seed(manualSeed)
        torch.cuda.manual_seed(manualSeed)
        torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False
        os.environ['PYTHONHASHSEED'] = str(manualSeed)

        # Root directory for dataset
        # dataroot = "data/celeba"

        # Number of workers for dataloader
        workers = 1

        # Batch size during training
        batch_size = 128

        # Spatial size of training images. All images will be resized to this
        # size using a transformer.
        #image_size = 64
        image_size = 32

        # Number of channels in the training images. For color images this is 3
        #nc = 3
        nc = 1

        # Size of z latent vector (i.e. size of generator input)
        nz = 100

        # Size of feature maps in generator
        #ngf = 64
        ngf = 8

        # Size of feature maps in discriminator
```

```
#ndf = 64
ndf = 8

# Number of training epochs
num_epochs = 5
num_epochs_wgan = 15
num_iters = 250

# Learning rate for optimizers
lr = 0.0002
lr_rms = 5e-4

# Beta1 hyperparam for Adam optimizers
beta1 = 0.5

# Number of GPUs available. Use 0 for CPU mode.
ngpu = 1

# Decide which device we want to run on
device = torch.device("cuda:0" if (torch.cuda.is_available() and ngpu > 0) else "cpu")

# Initialize BCELoss function
criterion = nn.BCELoss()

# Create batch of latent vectors that we will use to visualize
# the progression of the generator
fixed_noise = torch.randn(64, nz, 1, 1, device=device)

# Establish convention for real and fake labels during training
real_label = 1.0
fake_label = 0.0

# Several useful functions
def initialize_net(net_class, init_method, device, ngpu):

    # Create the generator
    net_inst = net_class(ngpu).to(device)

    # Handle multi-gpu if desired
    if (device.type == 'cuda') and (ngpu > 1):
        net_inst = nn.DataParallel(net_inst, list(range(ngpu)))

    # Apply the weights_init function to randomly initialize all weights
    # to mean=0, stdev=0.2.
    if init_method is not None:
        net_inst.apply(init_method)

    # Print the model
    print(net_inst)

    return net_inst

def plot_GAN_loss(losses, labels):

    plt.figure(figsize=(10,5))
    plt.title("Losses During Training")
```

```

for loss, label in zip(losses, labels):
    plt.plot(loss, label=f"{label}")

plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()

def plot_real_fake_images(real_batch, fake_batch):

    # Plot the real images
    plt.figure(figsize=(15,15))
    plt.subplot(1,2,1)
    plt.axis("off")
    plt.title("Real Images")
    plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=5, normalize=True).cpu(), (1,2,0)))

    # Plot the fake images from the last epoch
    plt.subplot(1,2,2)
    plt.axis("off")
    plt.title("Fake Images")
    plt.imshow(np.transpose(fake_batch[-1], (1,2,0)))
    plt.show()

# custom weights initialization called on netG and netD
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)

# Download the MNIST dataset
dataset = dset.MNIST(
    'data', train=True, download=True,
    transform=transforms.Compose([
        transforms.Resize(image_size), # Resize from 28 x 28 to 32 x 32 (so power of 2)
        transforms.CenterCrop(image_size),
        transforms.ToTensor(),
        transforms.Normalize((0.5,), (0.5,))
    ]))

# Create the dataloader
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                         shuffle=True, num_workers=workers)

# Plot some training images
real_batch = next(iter(dataloader))
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")

```

```
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64], padding=2, normalize=True).cpu(),(1,2,0)))
```

Random Seed: 999

Out[2]: <matplotlib.image.AxesImage at 0x18d47080588>



Architectural design for generator and discriminator

```

In [4]: # Generator Code
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution, state size. nz x 1 x 1
            nn.ConvTranspose2d( nz, ngf * 4, kernel_size=4, stride=1, padding=
0, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True), # inplace ReLU
            # current state size. (ngf*4) x 4 x 4
            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # current state size. (ngf*2) x 8 x 8
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # current state size. ngf x 16 x 16
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            # current state size. nc x 32 x 32
            # Produce number between -1 and 1, as pixel values have been norma
            lized to be between -1 and 1
            nn.Tanh()
        )

    def forward(self, input):
        return self.main(input)

class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 32 x 32
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 16 x 16
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 8 x 8
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 4 x 4
            nn.Conv2d(ndf * 4, 1, 4, 1, 0, bias=False),
            # state size. (ndf*4) x 1 x 1
            nn.Sigmoid() # Produce probability
        )

    def forward(self, input):
        return self.main(input)

```

Loss function and Training function

```

In [8]: # Initialize networks
netG = initialize_net(Generator, weights_init, device, ngpu)
netD = initialize_net(Discriminator, weights_init, device, ngpu)

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize  $\log(D(x)) + \log(1 - D(G(z)))$ 
        #####
        ## Train with all-real batch
        netD.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, device=device)
        # Forward pass real batch through D
        output = netD(real_cpu).view(-1)
        # Calculate loss on all-real batch
        errD_real = criterion(output, label)
        # Calculate gradients for D in backward pass
        errD_real.backward()
        D_x = output.mean().item()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD(fake.detach()).view(-1)
        # Calculate D's loss on the all-fake batch
        errD_fake = criterion(output, label)
        # Calculate the gradients for this batch
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        # Add the gradients from the all-real and all-fake batches
        errD = errD_real + errD_fake
        # Update D

```



```

optimizerD.step()

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG.step()

# Output training stats
if i % 50 == 0:
    print('%d/%d [%d/%d] \tLoss_D: %.4f \tLoss_G: %.4f \tD(x): %.4f \tD(G(z)): %.4f / %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1

```

```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): Tanh()
  )
)
Discriminator(
  (main): Sequential(
    (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (9): Sigmoid()
  )
)
Starting Training Loop...
[0/5][0/469] Loss_D: 1.4493 Loss_G: 0.7415 D(x): 0.4824 D(G(z)): 0.5060 / 0.4805
[0/5][50/469] Loss_D: 0.5649 Loss_G: 1.4601 D(x): 0.8009 D(G(z)): 0.2827 / 0.2385
[0/5][100/469] Loss_D: 0.3072 Loss_G: 2.2178 D(x): 0.8648 D(G(z)): 0.1438 / 0.1146
[0/5][150/469] Loss_D: 0.1532 Loss_G: 2.6514 D(x): 0.9455 D(G(z)): 0.0904 / 0.0784
[0/5][200/469] Loss_D: 0.0809 Loss_G: 3.3859 D(x): 0.9648 D(G(z)): 0.0432 / 0.0394
[0/5][250/469] Loss_D: 0.0421 Loss_G: 3.9898 D(x): 0.9821 D(G(z)): 0.0234 / 0.0215
[0/5][300/469] Loss_D: 0.0268 Loss_G: 4.3464 D(x): 0.9912 D(G(z)): 0.0177 / 0.0156
[0/5][350/469] Loss_D: 0.0232 Loss_G: 4.4663 D(x): 0.9930 D(G(z)): 0.01

```

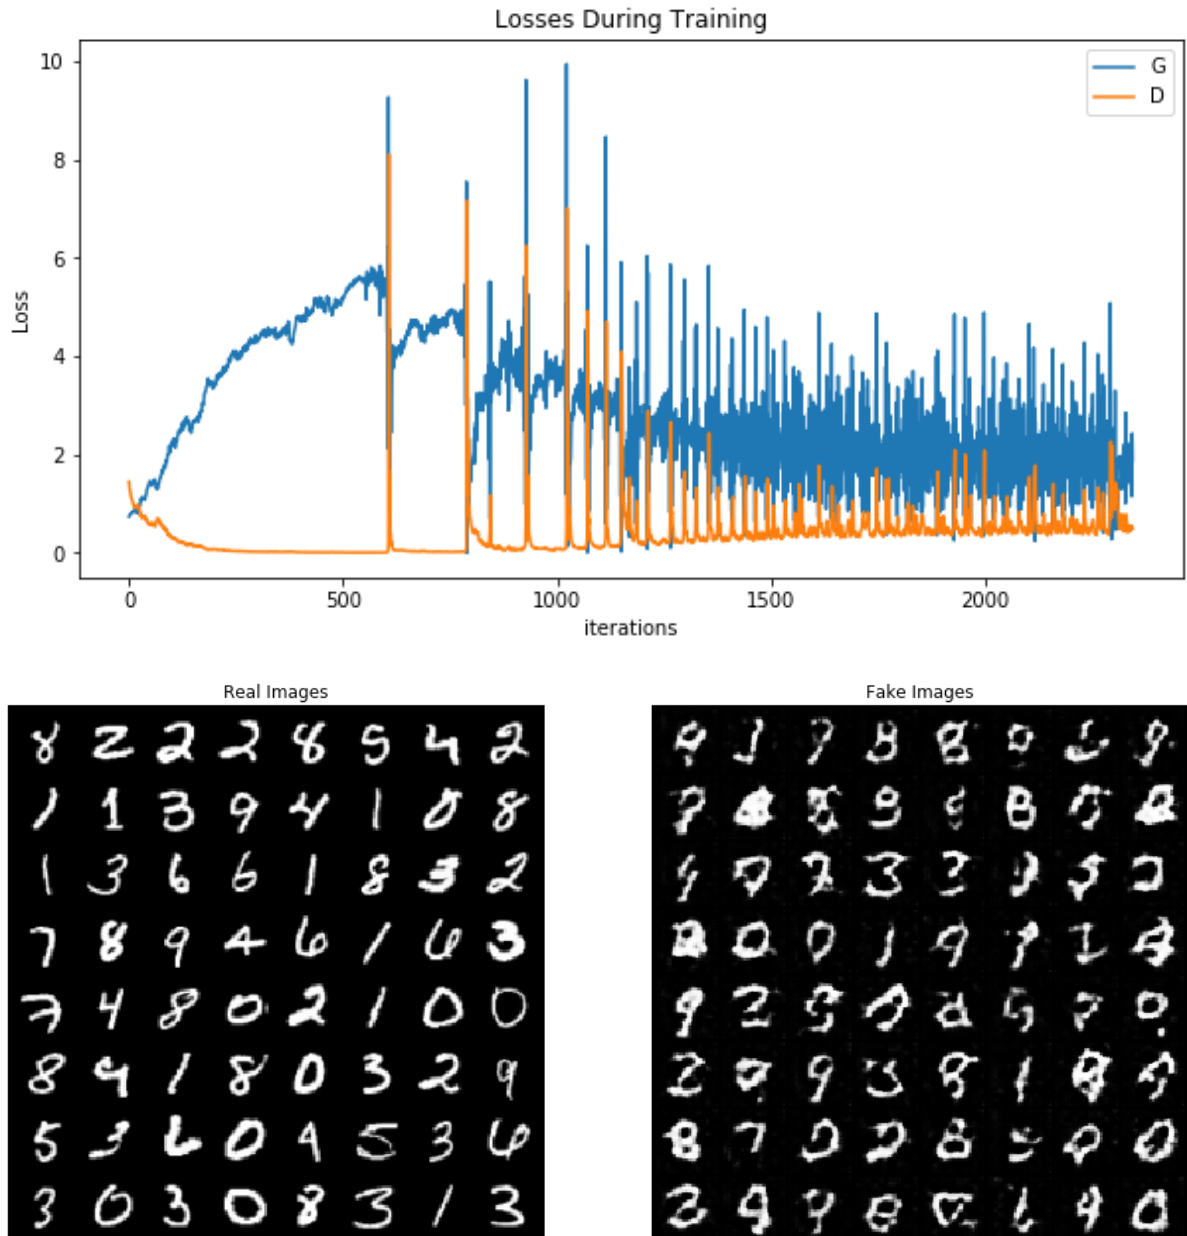
60 / 0.0138					
[0/5][400/469]	Loss_D: 0.0161	Loss_G: 4.7768	D(x): 0.9955	D(G(z)): 0.01	
15 / 0.0091					
[0/5][450/469]	Loss_D: 0.0099	Loss_G: 5.1826	D(x): 0.9978	D(G(z)): 0.00	
76 / 0.0066					
[1/5][0/469]	Loss_D: 0.0127	Loss_G: 5.0169	D(x): 0.9975	D(G(z)): 0.01	
01 / 0.0073					
[1/5][50/469]	Loss_D: 0.0095	Loss_G: 5.4539	D(x): 0.9966	D(G(z)): 0.00	
60 / 0.0050					
[1/5][100/469]	Loss_D: 0.0106	Loss_G: 5.5828	D(x): 0.9947	D(G(z)): 0.00	
52 / 0.0039					
[1/5][150/469]	Loss_D: 0.0707	Loss_G: 4.1313	D(x): 0.9751	D(G(z)): 0.04	
38 / 0.0174					
[1/5][200/469]	Loss_D: 0.0339	Loss_G: 4.6644	D(x): 0.9839	D(G(z)): 0.01	
73 / 0.0111					
[1/5][250/469]	Loss_D: 0.0243	Loss_G: 4.7487	D(x): 0.9870	D(G(z)): 0.01	
10 / 0.0097					
[1/5][300/469]	Loss_D: 0.0263	Loss_G: 4.6570	D(x): 0.9838	D(G(z)): 0.00	
99 / 0.0109					
[1/5][350/469]	Loss_D: 0.2402	Loss_G: 3.1571	D(x): 0.8952	D(G(z)): 0.11	
06 / 0.0513					
[1/5][400/469]	Loss_D: 0.1165	Loss_G: 3.5505	D(x): 0.9490	D(G(z)): 0.05	
72 / 0.0379					
[1/5][450/469]	Loss_D: 0.0839	Loss_G: 3.6389	D(x): 0.9443	D(G(z)): 0.01	
47 / 0.0361					
[2/5][0/469]	Loss_D: 0.2277	Loss_G: 2.5304	D(x): 0.9503	D(G(z)): 0.15	
79 / 0.0913					
[2/5][50/469]	Loss_D: 0.0605	Loss_G: 3.8779	D(x): 0.9763	D(G(z)): 0.03	
55 / 0.0236					
[2/5][100/469]	Loss_D: 0.1419	Loss_G: 3.1828	D(x): 0.9319	D(G(z)): 0.06	
70 / 0.0496					
[2/5][150/469]	Loss_D: 0.1373	Loss_G: 3.2217	D(x): 0.9511	D(G(z)): 0.08	
17 / 0.0465					
[2/5][200/469]	Loss_D: 0.1551	Loss_G: 2.8188	D(x): 0.9248	D(G(z)): 0.07	
05 / 0.0715					
[2/5][250/469]	Loss_D: 0.3293	Loss_G: 3.3079	D(x): 0.7501	D(G(z)): 0.02	
67 / 0.0410					
[2/5][300/469]	Loss_D: 0.1709	Loss_G: 2.7689	D(x): 0.9298	D(G(z)): 0.09	
01 / 0.0727					
[2/5][350/469]	Loss_D: 0.2094	Loss_G: 2.5953	D(x): 0.8997	D(G(z)): 0.09	
47 / 0.0850					
[2/5][400/469]	Loss_D: 0.2756	Loss_G: 2.4953	D(x): 0.8781	D(G(z)): 0.12	
89 / 0.0956					
[2/5][450/469]	Loss_D: 0.2437	Loss_G: 2.3766	D(x): 0.8823	D(G(z)): 0.10	
64 / 0.1074					
[3/5][0/469]	Loss_D: 0.5736	Loss_G: 0.8547	D(x): 0.6049	D(G(z)): 0.02	
86 / 0.4503					
[3/5][50/469]	Loss_D: 0.2849	Loss_G: 2.3658	D(x): 0.9202	D(G(z)): 0.17	
65 / 0.1052					
[3/5][100/469]	Loss_D: 0.7536	Loss_G: 2.7158	D(x): 0.9708	D(G(z)): 0.49	
35 / 0.0768					
[3/5][150/469]	Loss_D: 0.5502	Loss_G: 2.6759	D(x): 0.9472	D(G(z)): 0.36	
96 / 0.0825					
[3/5][200/469]	Loss_D: 0.3190	Loss_G: 2.1664	D(x): 0.8549	D(G(z)): 0.14	
02 / 0.1291					
[3/5][250/469]	Loss_D: 0.9409	Loss_G: 1.0806	D(x): 0.4383	D(G(z)): 0.03	
58 / 0.3602					

[3/5][300/469] 24 / 0.0953	Loss_D: 0.3255	Loss_G: 2.4825	D(x): 0.8707	D(G(z)): 0.16
[3/5][350/469] 72 / 0.1557	Loss_D: 0.4015	Loss_G: 2.0034	D(x): 0.8289	D(G(z)): 0.17
[3/5][400/469] 09 / 0.2106	Loss_D: 0.4051	Loss_G: 1.6725	D(x): 0.7982	D(G(z)): 0.15
[3/5][450/469] 85 / 0.1133	Loss_D: 0.5482	Loss_G: 2.3155	D(x): 0.8970	D(G(z)): 0.33
[4/5][0/469] 33 / 0.0993	Loss_D: 0.3806	Loss_G: 2.4335	D(x): 0.8588	D(G(z)): 0.19
[4/5][50/469] 05 / 0.0336	Loss_D: 0.5730	Loss_G: 3.5892	D(x): 0.9247	D(G(z)): 0.37
[4/5][100/469] 45 / 0.2139	Loss_D: 0.4288	Loss_G: 1.6665	D(x): 0.7756	D(G(z)): 0.14
[4/5][150/469] 64 / 0.3342	Loss_D: 0.4800	Loss_G: 1.1722	D(x): 0.7310	D(G(z)): 0.13
[4/5][200/469] 49 / 0.0510	Loss_D: 0.4490	Loss_G: 3.1399	D(x): 0.9111	D(G(z)): 0.28
[4/5][250/469] 90 / 0.3419	Loss_D: 0.5594	Loss_G: 1.1648	D(x): 0.6734	D(G(z)): 0.10
[4/5][300/469] 99 / 0.2102	Loss_D: 0.4193	Loss_G: 1.6752	D(x): 0.7690	D(G(z)): 0.12
[4/5][350/469] 95 / 0.0928	Loss_D: 0.4460	Loss_G: 2.5648	D(x): 0.8833	D(G(z)): 0.25
[4/5][400/469] 06 / 0.0376	Loss_D: 0.9967	Loss_G: 3.4140	D(x): 0.9776	D(G(z)): 0.60
[4/5][450/469] 39 / 0.1564	Loss_D: 0.5906	Loss_G: 1.9718	D(x): 0.7761	D(G(z)): 0.26

Visualization of the results

```
In [9]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



Task 1.1 Ablation study on batch normalization

1. Please modify the code provided in the Task 1.0 so that the neural network architecture does not contain any batch normalization layer.

Hint: modify the **Architectural design for generator and discriminator** section in Task 1.0

2. Train the model with modified networks and visualize the results.

```

In [5]: # Generator Code
class Generator_woBN(nn.Module):
    def __init__(self, ngpu):
        super(Generator_woBN, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            ##### YOUR CODE #####
            #####

            nn.ConvTranspose2d( nz, ngf * 4, kernel_size=4, stride=1, padding=
0, bias=False),
            ## nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True), # inplace ReLU
            # current state size. (ngf*4) x 4 x 4
            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.ReLU(True),
            # current state size. (ngf*2) x 8 x 8
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.ReLU(True),
            # current state size. ngf x 16 x 16
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            # current state size. nc x 32 x 32
            # Produce number between -1 and 1, as pixel values have been norma
lized to be between -1 and 1
            nn.Tanh()

            ##### END YOUR CODE #####
            #####
        )

    def forward(self, input):
        return self.main(input)

class Discriminator_woBN(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator_woBN, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            ##### YOUR CODE #####
            #####

            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 16 x 16
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 8 x 8
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 4 x 4
            nn.Conv2d(ndf * 4, 1, 4, 1, 0, bias=False),
            # state size. (ndf*4) x 1 x 1
            nn.Sigmoid() # Produce probability

```

```
##### END YOUR CODE #####
#####
    )
```

```
def forward(self, input):
    return self.main(input)
```

```
netG_noBN = initialize_net(Generator_woBN, weights_init, device, ngpu)
netD_noBN = initialize_net(Discriminator_woBN, weights_init, device, ngpu)
```

```
Generator_woBN(
    (main): Sequential(
      (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
      (1): ReLU(inplace=True)
      (2): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (3): ReLU(inplace=True)
      (4): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (5): ReLU(inplace=True)
      (6): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (7): Tanh()
    )
)
Discriminator_woBN(
    (main): Sequential(
      (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (1): LeakyReLU(negative_slope=0.2, inplace=True)
      (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (3): LeakyReLU(negative_slope=0.2, inplace=True)
      (4): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (5): LeakyReLU(negative_slope=0.2, inplace=True)
      (6): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
      (7): Sigmoid()
    )
)
```

```

In [22]: # Setup Adam optimizers for both G and D
optimizerD_noBN = optim.Adam(netD_noBN.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG_noBN = optim.Adam(netG_noBN.parameters(), lr=lr, betas=(beta1, 0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        ## Train with all-real batch
        netD_noBN.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, device=device)
        # Forward pass real batch through D
        output = netD_noBN(real_cpu).view(-1)
        # Calculate loss on all-real batch
        errD_real = criterion(output, label)
        # Calculate gradients for D in backward pass
        errD_real.backward()
        D_x = output.mean().item()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG_noBN(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD_noBN(fake.detach()).view(-1)
        # Calculate D's loss on the all-fake batch
        errD_fake = criterion(output, label)
        # Calculate the gradients for this batch
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        # Add the gradients from the all-real and all-fake batches
        errD = errD_real + errD_fake
        # Update D
        optimizerD_noBN.step()

        #####

```



```

# (2) Update G network: maximize Log(D(G(z)))
#####
netG_noBN.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD_noBN(fake).view(-1)
# Calculate G's Loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG_noBN.step()

# Output training stats
if i % 50 == 0:
    print('[%d/%d][%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.4f\tD(G(z)): %.4f / %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG_noBN(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1

```




Starting Training Loop...

[0/5][0/469]	Loss_D: 1.3862	Loss_G: 0.6931	D(x): 0.5001	D(G(z)): 0.50
00 / 0.5000				
[0/5][50/469]	Loss_D: 1.1021	Loss_G: 0.4735	D(x): 0.8969	D(G(z)): 0.62
92 / 0.6230				
[0/5][100/469]	Loss_D: 1.0617	Loss_G: 0.7312	D(x): 0.6825	D(G(z)): 0.49
29 / 0.4814				
[0/5][150/469]	Loss_D: 0.7274	Loss_G: 0.7138	D(x): 0.9526	D(G(z)): 0.49
09 / 0.4903				
[0/5][200/469]	Loss_D: 1.0531	Loss_G: 0.4995	D(x): 0.9002	D(G(z)): 0.60
77 / 0.6074				
[0/5][250/469]	Loss_D: 1.2455	Loss_G: 0.4615	D(x): 0.7923	D(G(z)): 0.63
47 / 0.6307				
[0/5][300/469]	Loss_D: 1.0370	Loss_G: 0.6426	D(x): 0.7651	D(G(z)): 0.53
15 / 0.5261				
[0/5][350/469]	Loss_D: 0.6507	Loss_G: 0.8827	D(x): 0.8804	D(G(z)): 0.39
78 / 0.4154				
[0/5][400/469]	Loss_D: 0.6759	Loss_G: 0.7737	D(x): 0.9547	D(G(z)): 0.46
46 / 0.4615				
[0/5][450/469]	Loss_D: 0.7132	Loss_G: 0.7028	D(x): 0.9813	D(G(z)): 0.49
96 / 0.4953				
[1/5][0/469]	Loss_D: 0.7360	Loss_G: 0.6875	D(x): 0.9709	D(G(z)): 0.50
41 / 0.5029				
[1/5][50/469]	Loss_D: 0.7618	Loss_G: 0.6754	D(x): 0.9633	D(G(z)): 0.50
85 / 0.5092				
[1/5][100/469]	Loss_D: 0.6918	Loss_G: 0.7167	D(x): 0.9872	D(G(z)): 0.49
19 / 0.4884				
[1/5][150/469]	Loss_D: 0.6848	Loss_G: 0.7120	D(x): 0.9922	D(G(z)): 0.49
11 / 0.4907				
[1/5][200/469]	Loss_D: 0.6991	Loss_G: 0.6916	D(x): 0.9965	D(G(z)): 0.50
12 / 0.5008				
[1/5][250/469]	Loss_D: 0.7248	Loss_G: 0.6830	D(x): 0.9852	D(G(z)): 0.50
61 / 0.5051				
[1/5][300/469]	Loss_D: 1.2294	Loss_G: 0.3710	D(x): 0.9858	D(G(z)): 0.70
20 / 0.6904				
[1/5][350/469]	Loss_D: 0.7619	Loss_G: 0.6635	D(x): 0.9703	D(G(z)): 0.51
81 / 0.5151				
[1/5][400/469]	Loss_D: 0.7359	Loss_G: 0.6849	D(x): 0.9716	D(G(z)): 0.50
57 / 0.5042				
[1/5][450/469]	Loss_D: 0.8494	Loss_G: 0.5832	D(x): 0.9751	D(G(z)): 0.56
08 / 0.5582				
[2/5][0/469]	Loss_D: 0.8059	Loss_G: 0.6273	D(x): 0.9642	D(G(z)): 0.53
62 / 0.5341				
[2/5][50/469]	Loss_D: 0.7459	Loss_G: 0.6861	D(x): 0.9661	D(G(z)): 0.50
78 / 0.5045				
[2/5][100/469]	Loss_D: 0.8877	Loss_G: 0.5797	D(x): 0.9443	D(G(z)): 0.56
32 / 0.5604				
[2/5][150/469]	Loss_D: 0.8267	Loss_G: 0.6318	D(x): 0.9409	D(G(z)): 0.53
30 / 0.5324				
[2/5][200/469]	Loss_D: 0.9887	Loss_G: 0.5047	D(x): 0.9676	D(G(z)): 0.61
46 / 0.6039				
[2/5][250/469]	Loss_D: 0.7888	Loss_G: 0.6859	D(x): 0.9247	D(G(z)): 0.50
71 / 0.5038				
[2/5][300/469]	Loss_D: 0.8040	Loss_G: 0.6379	D(x): 0.9626	D(G(z)): 0.53
45 / 0.5286				
[2/5][350/469]	Loss_D: 0.7322	Loss_G: 0.8264	D(x): 0.8649	D(G(z)): 0.43
33 / 0.4390				

[2/5][400/469] 96 / 0.5006	Loss_D: 0.7539	Loss_G: 0.6926	D(x): 0.9609	D(G(z)): 0.50
[2/5][450/469] 47 / 0.5309	Loss_D: 0.8252	Loss_G: 0.6348	D(x): 0.9456	D(G(z)): 0.53
[3/5][0/469] 89 / 0.4886	Loss_D: 0.7772	Loss_G: 0.7198	D(x): 0.9234	D(G(z)): 0.49
[3/5][50/469] 31 / 0.5647	Loss_D: 0.9471	Loss_G: 0.5730	D(x): 0.9151	D(G(z)): 0.57
[3/5][100/469] 96 / 0.5689	Loss_D: 0.9725	Loss_G: 0.5667	D(x): 0.9097	D(G(z)): 0.57
[3/5][150/469] 87 / 0.5733	Loss_D: 0.9877	Loss_G: 0.5592	D(x): 0.8919	D(G(z)): 0.57
[3/5][200/469] 40 / 0.5179	Loss_D: 0.8591	Loss_G: 0.6601	D(x): 0.8772	D(G(z)): 0.51
[3/5][250/469] 05 / 0.4616	Loss_D: 0.7949	Loss_G: 0.7761	D(x): 0.9287	D(G(z)): 0.51
[3/5][300/469] 64 / 0.5066	Loss_D: 0.8385	Loss_G: 0.6871	D(x): 0.9104	D(G(z)): 0.51
[3/5][350/469] 74 / 0.5442	Loss_D: 0.8649	Loss_G: 0.6124	D(x): 0.8698	D(G(z)): 0.50
[3/5][400/469] 83 / 0.5449	Loss_D: 0.9474	Loss_G: 0.6092	D(x): 0.9700	D(G(z)): 0.59
[3/5][450/469] 93 / 0.5443	Loss_D: 0.9090	Loss_G: 0.6099	D(x): 0.8803	D(G(z)): 0.53
[4/5][0/469] 19 / 0.4990	Loss_D: 0.8197	Loss_G: 0.6970	D(x): 0.9251	D(G(z)): 0.52
[4/5][50/469] 19 / 0.4910	Loss_D: 0.8413	Loss_G: 0.7134	D(x): 0.9088	D(G(z)): 0.52
[4/5][100/469] 70 / 0.4633	Loss_D: 0.7529	Loss_G: 0.7709	D(x): 0.9396	D(G(z)): 0.49
[4/5][150/469] 31 / 0.4791	Loss_D: 0.7445	Loss_G: 0.7366	D(x): 0.9086	D(G(z)): 0.47
[4/5][200/469] 90 / 0.5008	Loss_D: 0.7977	Loss_G: 0.6930	D(x): 0.9411	D(G(z)): 0.51
[4/5][250/469] 17 / 0.5021	Loss_D: 0.7729	Loss_G: 0.6904	D(x): 0.9492	D(G(z)): 0.51
[4/5][300/469] 37 / 0.5002	Loss_D: 0.7794	Loss_G: 0.6933	D(x): 0.9647	D(G(z)): 0.52
[4/5][350/469] 94 / 0.6126	Loss_D: 0.9193	Loss_G: 0.4914	D(x): 0.8252	D(G(z)): 0.50
[4/5][400/469] 36 / 0.5053	Loss_D: 0.7829	Loss_G: 0.6836	D(x): 0.9483	D(G(z)): 0.51
[4/5][450/469] 94 / 0.4766	Loss_D: 0.7414	Loss_G: 0.7421	D(x): 0.9389	D(G(z)): 0.48

```
In [23]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



Task 1.2 Ablation study on the trick: "Construct different mini-batches for real and fake"

1. Please modify the code provided in the Task 1.0 so that the discriminator algorithm part computes the forward and backward pass for fake and real images concatenated together (with their corresponding fake and real labels concatenated as well) instead of computing the forward and backward passes for fake and real images separately.

Hint: modify the ***Loss function and Training function*** section in Task 1.0.

2. Train the model with modified networks and visualize the results.

```

In [55]: # re-initilize networks for the generator and discriminator.
netG = initialize_net(Generator, weights_init, device, ngpu)
netD = initialize_net(Discriminator, weights_init, device, ngpu)

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize  $\log(D(x)) + \log(1 - D(G(z)))$ 
        #####

        ##### YOUR CODE #####

        #####

        netD.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)

        noise = torch.randn(b_size, nz, 1, 1, device=device)
        fake = netG(noise)

        data_concat = torch.cat((real_cpu, fake))
        label_concat = torch.cat((torch.full((b_size,), real_label, device=device),
                                torch.full((b_size,), fake_label, device=device)))

        output_concat = netD(data_concat).view(-1)

        errD = criterion(output_concat, label_concat)
        errD.backward()

        # Update D
        optimizerD.step()

        ##### END YOUR CODE #####

    ###

```

```

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
noise = torch.randn(b_size, nz, 1, 1, device=device)
fake = netG(noise)

label = torch.full((b_size,), real_label, device=device) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG.step()

# Output training stats
if i % 50 == 0:
    print('%d/%d [%d/%d] \t Loss_D: %.4f \t Loss_G: %.4f \t D(G(z)): %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_G_z2))

# Save losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1

```

```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): Tanh()
  )
)
Discriminator(
  (main): Sequential(
    (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (9): Sigmoid()
  )
)

```

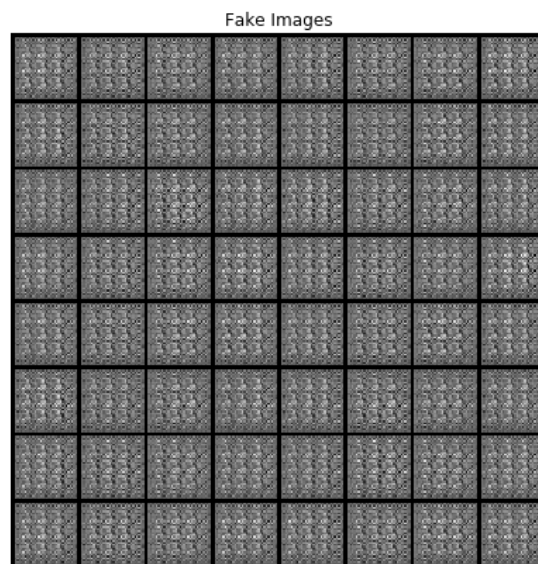

Starting Training Loop...

[0/5][0/469]	Loss_D: 0.7281	Loss_G: 0.7642	D(G(z)): 0.4703
[0/5][50/469]	Loss_D: 0.0334	Loss_G: 0.1177	D(G(z)): 0.8895
[0/5][100/469]	Loss_D: 0.0121	Loss_G: 0.0459	D(G(z)): 0.9552
[0/5][150/469]	Loss_D: 0.0066	Loss_G: 0.0263	D(G(z)): 0.9740
[0/5][200/469]	Loss_D: 0.0043	Loss_G: 0.0182	D(G(z)): 0.9820
[0/5][250/469]	Loss_D: 0.0030	Loss_G: 0.0128	D(G(z)): 0.9873
[0/5][300/469]	Loss_D: 0.0022	Loss_G: 0.0095	D(G(z)): 0.9905
[0/5][350/469]	Loss_D: 0.0018	Loss_G: 0.0075	D(G(z)): 0.9925
[0/5][400/469]	Loss_D: 0.0014	Loss_G: 0.0062	D(G(z)): 0.9938
[0/5][450/469]	Loss_D: 0.0012	Loss_G: 0.0051	D(G(z)): 0.9950
[1/5][0/469]	Loss_D: 0.0010	Loss_G: 0.0048	D(G(z)): 0.9952
[1/5][50/469]	Loss_D: 0.0009	Loss_G: 0.0041	D(G(z)): 0.9959
[1/5][100/469]	Loss_D: 0.0008	Loss_G: 0.0035	D(G(z)): 0.9965
[1/5][150/469]	Loss_D: 0.0006	Loss_G: 0.0031	D(G(z)): 0.9969
[1/5][200/469]	Loss_D: 0.0006	Loss_G: 0.0028	D(G(z)): 0.9972

[1/5][250/469]	Loss_D: 0.0005	Loss_G: 0.0025	D(G(z)): 0.9975
[1/5][300/469]	Loss_D: 0.0004	Loss_G: 0.0022	D(G(z)): 0.9978
[1/5][350/469]	Loss_D: 0.0004	Loss_G: 0.0020	D(G(z)): 0.9980
[1/5][400/469]	Loss_D: 0.0004	Loss_G: 0.0018	D(G(z)): 0.9982
[1/5][450/469]	Loss_D: 0.0003	Loss_G: 0.0017	D(G(z)): 0.9983
[2/5][0/469]	Loss_D: 0.0003	Loss_G: 0.0017	D(G(z)): 0.9983
[2/5][50/469]	Loss_D: 0.0003	Loss_G: 0.0015	D(G(z)): 0.9985
[2/5][100/469]	Loss_D: 0.0003	Loss_G: 0.0014	D(G(z)): 0.9986
[2/5][150/469]	Loss_D: 0.0002	Loss_G: 0.0013	D(G(z)): 0.9987
[2/5][200/469]	Loss_D: 0.0002	Loss_G: 0.0012	D(G(z)): 0.9988
[2/5][250/469]	Loss_D: 0.0002	Loss_G: 0.0011	D(G(z)): 0.9989
[2/5][300/469]	Loss_D: 0.0002	Loss_G: 0.0011	D(G(z)): 0.9989
[2/5][350/469]	Loss_D: 0.0002	Loss_G: 0.0010	D(G(z)): 0.9990
[2/5][400/469]	Loss_D: 0.0002	Loss_G: 0.0009	D(G(z)): 0.9991
[2/5][450/469]	Loss_D: 0.0002	Loss_G: 0.0009	D(G(z)): 0.9991
[3/5][0/469]	Loss_D: 0.0001	Loss_G: 0.0008	D(G(z)): 0.9992
[3/5][50/469]	Loss_D: 0.0001	Loss_G: 0.0008	D(G(z)): 0.9992
[3/5][100/469]	Loss_D: 0.0001	Loss_G: 0.0007	D(G(z)): 0.9993
[3/5][150/469]	Loss_D: 0.0001	Loss_G: 0.0007	D(G(z)): 0.9993
[3/5][200/469]	Loss_D: 0.0001	Loss_G: 0.0007	D(G(z)): 0.9993
[3/5][250/469]	Loss_D: 0.0001	Loss_G: 0.0006	D(G(z)): 0.9994
[3/5][300/469]	Loss_D: 0.0001	Loss_G: 0.0006	D(G(z)): 0.9994
[3/5][350/469]	Loss_D: 0.0001	Loss_G: 0.0006	D(G(z)): 0.9994
[3/5][400/469]	Loss_D: 0.0001	Loss_G: 0.0005	D(G(z)): 0.9995
[3/5][450/469]	Loss_D: 0.0001	Loss_G: 0.0005	D(G(z)): 0.9995
[4/5][0/469]	Loss_D: 0.0001	Loss_G: 0.0005	D(G(z)): 0.9995
[4/5][50/469]	Loss_D: 0.0001	Loss_G: 0.0005	D(G(z)): 0.9995
[4/5][100/469]	Loss_D: 0.0001	Loss_G: 0.0005	D(G(z)): 0.9995
[4/5][150/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][200/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][250/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][300/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][350/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][400/469]	Loss_D: 0.0001	Loss_G: 0.0004	D(G(z)): 0.9996
[4/5][450/469]	Loss_D: 0.0001	Loss_G: 0.0003	D(G(z)): 0.9997

```
In [56]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



Task 1.3 Ablation study on the generator's loss function

1. Please modify the code provided in the Task 1.0 so that the *Generator* algorithm part minimizes $\log(1 - D(G(z)))$ instead of the modified loss function suggested in the original GAN paper of $-\log(D(G(z)))$.
 - A. Modify the **Loss function and Training function** section in Task 1.0
 - B. (Hint) Try to understand the definition of [BCE loss](https://pytorch.org/docs/stable/generated/torch.nn.BCELoss.html) (<https://pytorch.org/docs/stable/generated/torch.nn.BCELoss.html>) first and how the modified loss function was implemented.
2. Train the model with modified networks and visualize the results.

```

In [12]: # re-initilize networks for the generator and discriminator.
netG = initialize_net(Generator, weights_init, device, ngpu)
netD = initialize_net(Discriminator, weights_init, device, ngpu)

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize  $\log(D(x)) + \log(1 - D(G(z)))$ 
        #####
        ## Train with all-real batch
        netD.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, device=device)
        # Forward pass real batch through D
        output = netD(real_cpu).view(-1)
        # Calculate loss on all-real batch
        errD_real = criterion(output, label)
        # Calculate gradients for D in backward pass
        errD_real.backward()
        D_x = output.mean().item()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD(fake.detach()).view(-1)
        # Calculate D's loss on the all-fake batch
        errD_fake = criterion(output, label)
        # Calculate the gradients for this batch
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        # Add the gradients from the all-real and all-fake batches
        errD = errD_real + errD_fake
        # Update D
        optimizerD.step()

```

```

#####
# (2) Update G network
#####

##### YOUR CODE #####
####
    ....
    #label.fill_(fake_label)
    # Classify all fake batch with D
    output = netD(fake).view(-1)
    # Calculate D's loss on the all-fake batch
    #errG = torch.log(1- output)
    #output = torch.log(1-output)
    errG = -criterion(output, label)
    # Calculate the gradients for this batch
    errG.backward()
    D_G_z2 = output.mean().item()
    # Update G
    optimizerG.step()
    ...

    netG.zero_grad()
    label.fill_(fake_label) # fake labels to minimize
    # Since we just updated D, perform another forward pass of all-fake batch through D
    output = netD(fake).view(-1)
    # Calculate G's loss based on this output
    errG = -criterion(output, label)
    # Calculate gradients for G
    errG.backward()
    D_G_z2 = output.mean().item()
    # Update G
    optimizerG.step()

##### END YOUR CODE #####
###

# Output training stats
if i % 50 == 0:
    print('[%d/%d][%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.4f\tD(G(z)): %.4f / %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

    iters += 1

```



```

Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): Tanh()
  )
)
Discriminator(
  (main): Sequential(
    (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (9): Sigmoid()
  )
)
Starting Training Loop...
[0/5][0/469] Loss_D: 1.4669 Loss_G: -0.6471 D(x): 0.4565 D(G(z)): 0.4869 / 0.4727
[0/5][50/469] Loss_D: 0.3643 Loss_G: -0.2051 D(x): 0.8811 D(G(z)): 0.2080 / 0.1839
[0/5][100/469] Loss_D: 0.3241 Loss_G: -0.1450 D(x): 0.8609 D(G(z)): 0.1497 / 0.1331
[0/5][150/469] Loss_D: 0.0572 Loss_G: -0.0282 D(x): 0.9748 D(G(z)): 0.0307 / 0.0277
[0/5][200/469] Loss_D: 0.0267 Loss_G: -0.0140 D(x): 0.9882 D(G(z)): 0.0147 / 0.0138
[0/5][250/469] Loss_D: 0.0162 Loss_G: -0.0088 D(x): 0.9929 D(G(z)): 0.0090 / 0.0087
[0/5][300/469] Loss_D: 0.0106 Loss_G: -0.0055 D(x): 0.9950 D(G(z)): 0.0055 / 0.0055
[0/5][350/469] Loss_D: 0.0086 Loss_G: -0.0047 D(x): 0.9963 D(G(z)): 0.00

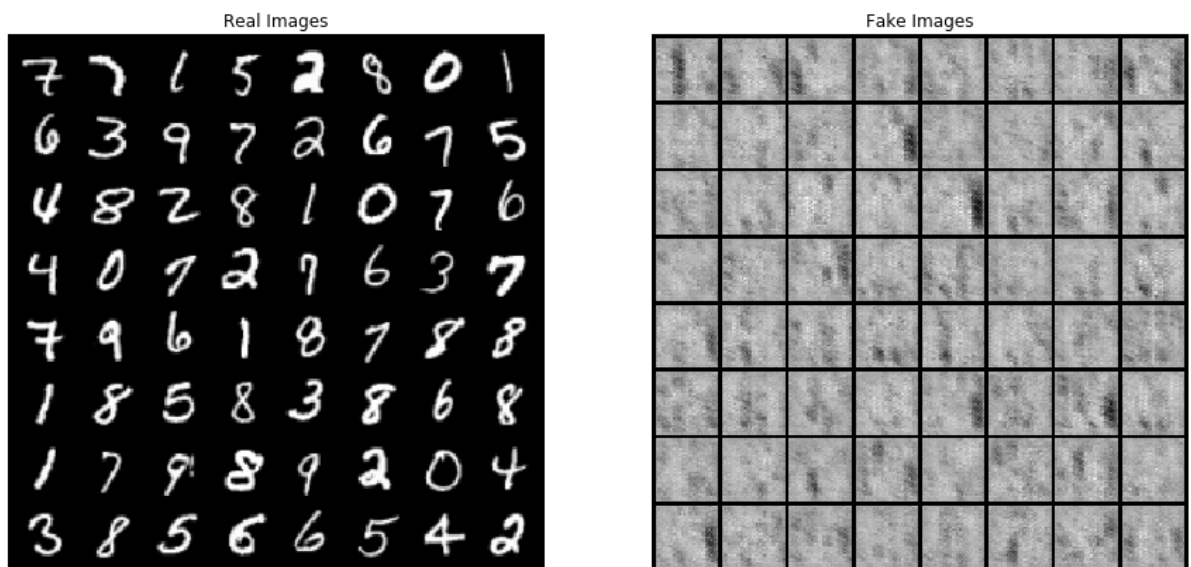

```

48 / 0.0047				
[0/5][400/469]	Loss_D: 0.0067	Loss_G: -0.0036	D(x): 0.9969	D(G(z)): 0.00
36 / 0.0036				
[0/5][450/469]	Loss_D: 0.0046	Loss_G: -0.0028	D(x): 0.9982	D(G(z)): 0.00
29 / 0.0028				
[1/5][0/469]	Loss_D: 0.0045	Loss_G: -0.0027	D(x): 0.9982	D(G(z)): 0.00
27 / 0.0026				
[1/5][50/469]	Loss_D: 0.0041	Loss_G: -0.0022	D(x): 0.9981	D(G(z)): 0.00
22 / 0.0022				
[1/5][100/469]	Loss_D: 0.0029	Loss_G: -0.0018	D(x): 0.9989	D(G(z)): 0.00
18 / 0.0018				
[1/5][150/469]	Loss_D: 0.0026	Loss_G: -0.0016	D(x): 0.9989	D(G(z)): 0.00
16 / 0.0016				
[1/5][200/469]	Loss_D: 0.0025	Loss_G: -0.0014	D(x): 0.9990	D(G(z)): 0.00
14 / 0.0014				
[1/5][250/469]	Loss_D: 0.0019	Loss_G: -0.0012	D(x): 0.9992	D(G(z)): 0.00
12 / 0.0012				
[1/5][300/469]	Loss_D: 0.0017	Loss_G: -0.0011	D(x): 0.9994	D(G(z)): 0.00
11 / 0.0011				
[1/5][350/469]	Loss_D: 0.0015	Loss_G: -0.0008	D(x): 0.9994	D(G(z)): 0.00
08 / 0.0008				
[1/5][400/469]	Loss_D: 0.0013	Loss_G: -0.0007	D(x): 0.9994	D(G(z)): 0.00
07 / 0.0007				
[1/5][450/469]	Loss_D: 0.0013	Loss_G: -0.0008	D(x): 0.9995	D(G(z)): 0.00
08 / 0.0008				
[2/5][0/469]	Loss_D: 0.0012	Loss_G: -0.0007	D(x): 0.9996	D(G(z)): 0.00
07 / 0.0007				
[2/5][50/469]	Loss_D: 0.0010	Loss_G: -0.0006	D(x): 0.9996	D(G(z)): 0.00
06 / 0.0006				
[2/5][100/469]	Loss_D: 0.0010	Loss_G: -0.0006	D(x): 0.9996	D(G(z)): 0.00
06 / 0.0006				
[2/5][150/469]	Loss_D: 0.0009	Loss_G: -0.0005	D(x): 0.9996	D(G(z)): 0.00
05 / 0.0005				
[2/5][200/469]	Loss_D: 0.0009	Loss_G: -0.0005	D(x): 0.9996	D(G(z)): 0.00
05 / 0.0005				
[2/5][250/469]	Loss_D: 0.0008	Loss_G: -0.0004	D(x): 0.9997	D(G(z)): 0.00
04 / 0.0004				
[2/5][300/469]	Loss_D: 0.0007	Loss_G: -0.0004	D(x): 0.9997	D(G(z)): 0.00
04 / 0.0004				
[2/5][350/469]	Loss_D: 0.0007	Loss_G: -0.0004	D(x): 0.9997	D(G(z)): 0.00
04 / 0.0004				
[2/5][400/469]	Loss_D: 0.0007	Loss_G: -0.0004	D(x): 0.9997	D(G(z)): 0.00
04 / 0.0004				
[2/5][450/469]	Loss_D: 0.0006	Loss_G: -0.0004	D(x): 0.9998	D(G(z)): 0.00
04 / 0.0004				
[3/5][0/469]	Loss_D: 0.0007	Loss_G: -0.0004	D(x): 0.9997	D(G(z)): 0.00
04 / 0.0004				
[3/5][50/469]	Loss_D: 0.0005	Loss_G: -0.0003	D(x): 0.9998	D(G(z)): 0.00
03 / 0.0003				
[3/5][100/469]	Loss_D: 0.0005	Loss_G: -0.0003	D(x): 0.9997	D(G(z)): 0.00
03 / 0.0003				
[3/5][150/469]	Loss_D: 0.0006	Loss_G: -0.0003	D(x): 0.9998	D(G(z)): 0.00
03 / 0.0003				
[3/5][200/469]	Loss_D: 0.0005	Loss_G: -0.0003	D(x): 0.9998	D(G(z)): 0.00
03 / 0.0003				
[3/5][250/469]	Loss_D: 0.0005	Loss_G: -0.0003	D(x): 0.9998	D(G(z)): 0.00
03 / 0.0003				

[3/5][300/469] 02 / 0.0002	Loss_D: 0.0004	Loss_G: -0.0002	D(x): 0.9998	D(G(z)): 0.00
[3/5][350/469] 02 / 0.0002	Loss_D: 0.0004	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[3/5][400/469] 02 / 0.0002	Loss_D: 0.0004	Loss_G: -0.0002	D(x): 0.9998	D(G(z)): 0.00
[3/5][450/469] 02 / 0.0002	Loss_D: 0.0004	Loss_G: -0.0002	D(x): 0.9998	D(G(z)): 0.00
[4/5][0/469] 02 / 0.0002	Loss_D: 0.0004	Loss_G: -0.0002	D(x): 0.9998	D(G(z)): 0.00
[4/5][50/469] 02 / 0.0002	Loss_D: 0.0003	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[4/5][100/469] 02 / 0.0002	Loss_D: 0.0003	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[4/5][150/469] 02 / 0.0002	Loss_D: 0.0003	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[4/5][200/469] 02 / 0.0002	Loss_D: 0.0003	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[4/5][250/469] 02 / 0.0002	Loss_D: 0.0003	Loss_G: -0.0002	D(x): 0.9999	D(G(z)): 0.00
[4/5][300/469] 01 / 0.0001	Loss_D: 0.0003	Loss_G: -0.0001	D(x): 0.9999	D(G(z)): 0.00
[4/5][350/469] 01 / 0.0001	Loss_D: 0.0002	Loss_G: -0.0001	D(x): 0.9999	D(G(z)): 0.00
[4/5][400/469] 01 / 0.0001	Loss_D: 0.0002	Loss_G: -0.0001	D(x): 0.9999	D(G(z)): 0.00
[4/5][450/469] 01 / 0.0001	Loss_D: 0.0002	Loss_G: -0.0001	D(x): 0.9999	D(G(z)): 0.00

```
In [13]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



In []:

Task 1.4 Ablation study on the weight initialization

1. Please use the function `initialize_net` provided in Task 1.0 to initialize the generator and discriminator function without weight initialization (HINT: There is no need to modify the code for `initialize_net` function).
2. Train the model with modified networks and visualize the results.

```
In [18]: ##### YOUR CODE #####
netD_woinit = initialize_net(Discriminator, None, device, ngpu)
netG_woinit = initialize_net(Generator, None, device, ngpu)
##### END YOUR CODE #####
```

```
Discriminator(
  (main): Sequential(
    (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (9): Sigmoid()
  )
)
Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): Tanh()
  )
)
```

```

In [19]: # Setup Adam optimizers for both G and D
optimizerD_woinit = optim.Adam(netD_woinit.parameters(), lr=lr, betas=(beta1,
0.999))
optimizerG_woinit = optim.Adam(netG_woinit.parameters(), lr=lr, betas=(beta1,
0.999))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        ## Train with all-real batch
        netD_woinit.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, device=device)
        # Forward pass real batch through D
        output = netD_woinit(real_cpu).view(-1)
        # Calculate loss on all-real batch
        errD_real = criterion(output, label)
        # Calculate gradients for D in backward pass
        errD_real.backward()
        D_x = output.mean().item()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG_woinit(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD_woinit(fake.detach()).view(-1)
        # Calculate D's loss on the all-fake batch
        errD_fake = criterion(output, label)
        # Calculate the gradients for this batch
        errD_fake.backward()
        D_G_z1 = output.mean().item()
        # Add the gradients from the all-real and all-fake batches
        errD = errD_real + errD_fake
        # Update D
        optimizerD_woinit.step()

        #####

```

```

# (2) Update G network: maximize Log(D(G(z)))
#####
netG_woinit.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch through D
output = netD_woinit(fake).view(-1)
# Calculate G's Loss based on this output
errG = criterion(output, label)
# Calculate gradients for G
errG.backward()
D_G_z2 = output.mean().item()
# Update G
optimizerG_woinit.step()

# Output training stats
if i % 50 == 0:
    print('[%d/%d][%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.4f\tD(G(z)): %.4f / %.4f'
          % (epoch, num_epochs, i, len(dataloader),
             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
    with torch.no_grad():
        fake = netG_woinit(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

iters += 1

```

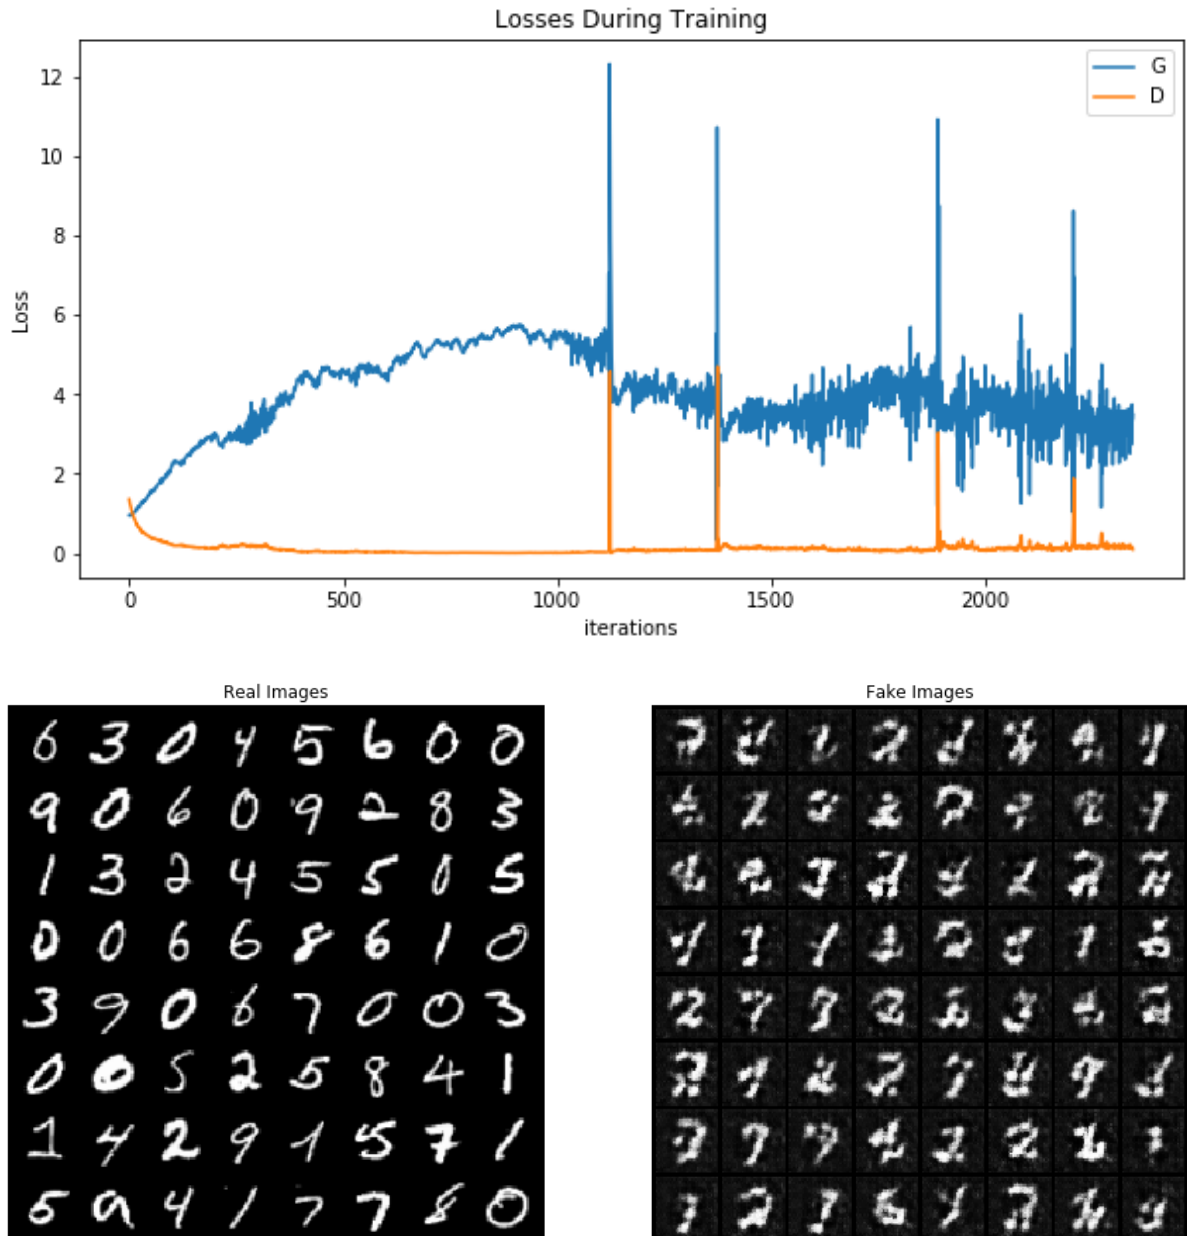
Starting Training Loop...

[0/5][0/469]	Loss_D: 1.3569	Loss_G: 0.9581	D(x): 0.4424	D(G(z)): 0.40
16 / 0.3917				
[0/5][50/469]	Loss_D: 0.3826	Loss_G: 1.5581	D(x): 0.8935	D(G(z)): 0.23
40 / 0.2161				
[0/5][100/469]	Loss_D: 0.2050	Loss_G: 2.2244	D(x): 0.9373	D(G(z)): 0.12
92 / 0.1130				
[0/5][150/469]	Loss_D: 0.1600	Loss_G: 2.5938	D(x): 0.9438	D(G(z)): 0.09
49 / 0.0789				
[0/5][200/469]	Loss_D: 0.1197	Loss_G: 2.9223	D(x): 0.9609	D(G(z)): 0.07
43 / 0.0589				
[0/5][250/469]	Loss_D: 0.1929	Loss_G: 2.9358	D(x): 0.9060	D(G(z)): 0.07
77 / 0.0605				
[0/5][300/469]	Loss_D: 0.1536	Loss_G: 3.3917	D(x): 0.9219	D(G(z)): 0.06
48 / 0.0386				
[0/5][350/469]	Loss_D: 0.0882	Loss_G: 3.8163	D(x): 0.9542	D(G(z)): 0.03
71 / 0.0247				
[0/5][400/469]	Loss_D: 0.0380	Loss_G: 4.3452	D(x): 0.9835	D(G(z)): 0.02
06 / 0.0145				
[0/5][450/469]	Loss_D: 0.0389	Loss_G: 4.6153	D(x): 0.9810	D(G(z)): 0.01
90 / 0.0118				
[1/5][0/469]	Loss_D: 0.0354	Loss_G: 4.5837	D(x): 0.9810	D(G(z)): 0.01
55 / 0.0120				
[1/5][50/469]	Loss_D: 0.0400	Loss_G: 4.3595	D(x): 0.9883	D(G(z)): 0.02
72 / 0.0151				
[1/5][100/469]	Loss_D: 0.0266	Loss_G: 4.5798	D(x): 0.9892	D(G(z)): 0.01
51 / 0.0111				
[1/5][150/469]	Loss_D: 0.0378	Loss_G: 4.7847	D(x): 0.9829	D(G(z)): 0.01
41 / 0.0094				
[1/5][200/469]	Loss_D: 0.0119	Loss_G: 5.2471	D(x): 0.9954	D(G(z)): 0.00
73 / 0.0057				
[1/5][250/469]	Loss_D: 0.0100	Loss_G: 5.3343	D(x): 0.9966	D(G(z)): 0.00
65 / 0.0052				
[1/5][300/469]	Loss_D: 0.0128	Loss_G: 5.3213	D(x): 0.9951	D(G(z)): 0.00
78 / 0.0052				
[1/5][350/469]	Loss_D: 0.0130	Loss_G: 5.4676	D(x): 0.9946	D(G(z)): 0.00
75 / 0.0047				
[1/5][400/469]	Loss_D: 0.0077	Loss_G: 5.5294	D(x): 0.9973	D(G(z)): 0.00
50 / 0.0042				
[1/5][450/469]	Loss_D: 0.0098	Loss_G: 5.6918	D(x): 0.9951	D(G(z)): 0.00
48 / 0.0036				
[2/5][0/469]	Loss_D: 0.0114	Loss_G: 5.2923	D(x): 0.9968	D(G(z)): 0.00
82 / 0.0055				
[2/5][50/469]	Loss_D: 0.0093	Loss_G: 5.4973	D(x): 0.9968	D(G(z)): 0.00
61 / 0.0043				
[2/5][100/469]	Loss_D: 0.0143	Loss_G: 5.3100	D(x): 0.9947	D(G(z)): 0.00
89 / 0.0054				
[2/5][150/469]	Loss_D: 0.0342	Loss_G: 4.9135	D(x): 0.9807	D(G(z)): 0.01
42 / 0.0082				
[2/5][200/469]	Loss_D: 0.0759	Loss_G: 3.8271	D(x): 0.9733	D(G(z)): 0.04
61 / 0.0253				
[2/5][250/469]	Loss_D: 0.0886	Loss_G: 3.9323	D(x): 0.9503	D(G(z)): 0.03
05 / 0.0257				
[2/5][300/469]	Loss_D: 0.0675	Loss_G: 4.1724	D(x): 0.9698	D(G(z)): 0.03
44 / 0.0209				
[2/5][350/469]	Loss_D: 0.1036	Loss_G: 3.6430	D(x): 0.9491	D(G(z)): 0.04
26 / 0.0344				

[2/5][400/469] 59 / 0.0276	Loss_D: 0.0748	Loss_G: 3.8228	D(x): 0.9742	D(G(z)): 0.04
[2/5][450/469] 32 / 0.0553	Loss_D: 0.1798	Loss_G: 3.0636	D(x): 0.9369	D(G(z)): 0.10
[3/5][0/469] 64 / 0.0508	Loss_D: 0.1467	Loss_G: 3.2363	D(x): 0.9337	D(G(z)): 0.06
[3/5][50/469] 95 / 0.0311	Loss_D: 0.1483	Loss_G: 3.8654	D(x): 0.9357	D(G(z)): 0.06
[3/5][100/469] 80 / 0.0585	Loss_D: 0.1401	Loss_G: 3.0381	D(x): 0.9309	D(G(z)): 0.05
[3/5][150/469] 72 / 0.0198	Loss_D: 0.0910	Loss_G: 4.2358	D(x): 0.9710	D(G(z)): 0.05
[3/5][200/469] 70 / 0.0831	Loss_D: 0.1488	Loss_G: 2.8348	D(x): 0.9139	D(G(z)): 0.04
[3/5][250/469] 77 / 0.0701	Loss_D: 0.1099	Loss_G: 2.9160	D(x): 0.9212	D(G(z)): 0.01
[3/5][300/469] 80 / 0.0259	Loss_D: 0.0954	Loss_G: 3.9063	D(x): 0.9425	D(G(z)): 0.02
[3/5][350/469] 06 / 0.0200	Loss_D: 0.0597	Loss_G: 4.2697	D(x): 0.9834	D(G(z)): 0.04
[3/5][400/469] 72 / 0.0235	Loss_D: 0.0492	Loss_G: 4.0862	D(x): 0.9697	D(G(z)): 0.01
[3/5][450/469] 98 / 0.0286	Loss_D: 0.0775	Loss_G: 3.9485	D(x): 0.9674	D(G(z)): 0.03
[4/5][0/469] 75 / 0.0355	Loss_D: 0.0941	Loss_G: 3.6881	D(x): 0.9330	D(G(z)): 0.01
[4/5][50/469] 19 / 0.0506	Loss_D: 0.0999	Loss_G: 3.4199	D(x): 0.9411	D(G(z)): 0.03
[4/5][100/469] 05 / 0.0482	Loss_D: 0.1172	Loss_G: 3.4347	D(x): 0.9255	D(G(z)): 0.03
[4/5][150/469] 35 / 0.0337	Loss_D: 0.0930	Loss_G: 3.8003	D(x): 0.9764	D(G(z)): 0.06
[4/5][200/469] 29 / 0.0293	Loss_D: 0.0839	Loss_G: 3.9679	D(x): 0.9536	D(G(z)): 0.03
[4/5][250/469] 09 / 0.0257	Loss_D: 0.1436	Loss_G: 4.0866	D(x): 0.9600	D(G(z)): 0.09
[4/5][300/469] 19 / 0.0584	Loss_D: 0.0960	Loss_G: 3.1791	D(x): 0.9422	D(G(z)): 0.03
[4/5][350/469] 46 / 0.0290	Loss_D: 0.1126	Loss_G: 3.9329	D(x): 0.9398	D(G(z)): 0.04
[4/5][400/469] 23 / 0.0933	Loss_D: 0.2129	Loss_G: 2.6459	D(x): 0.9199	D(G(z)): 0.11
[4/5][450/469] 10 / 0.1248	Loss_D: 0.1977	Loss_G: 2.3631	D(x): 0.8789	D(G(z)): 0.04

```
In [20]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



Exercise 2: Implement the WGAN with weight clipping

Wasserstein GAN ([WGAN \(https://arxiv.org/abs/1701.07875\)](https://arxiv.org/abs/1701.07875)) is an alternative training strategy to traditional GAN. WGAN may provide more stable learning and may avoid problems faced in traditional GAN training like mode collapse.

1. Rewrite the loss functions and training function according to the algorithm introduced in slide 18 in [Lecture note for WGAN](https://www.davidinouye.com/course/ece57000-fall-2021/lectures/wasserstein-gan.pdf) (<https://www.davidinouye.com/course/ece57000-fall-2021/lectures/wasserstein-gan.pdf>). A few notes/hints:
 - A. Keep the same generator as in Exercise 1, Task 1.0, but modify the discriminator so that there is no restriction on the range of the output. (Simply comment out the last `Sigmoid` layer)
 - B. Modify the optimizer to be the RMSProp optimizer with a learning rate equal to the value in `lr_rms` (which we set to $5e-4$, which is larger than the rate in the paper but works better for our purposes).
 - C. Use `torch.Tensor.clamp_()`. (https://pytorch.org/docs/stable/generated/torch.Tensor.clamp_.html) function to clip the parameter values. You will need to do this for all parameters of the discriminator. See algorithm for when to do this.
2. Train the model with modified networks and visualize the results.

```

In [6]: # Generator Code
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution, state size. nz x 1 x 1
            nn.ConvTranspose2d( nz, ngf * 4, kernel_size=4, stride=1, padding=
0, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True), # inplace ReLU
            # current state size. (ngf*4) x 4 x 4
            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # current state size. (ngf*2) x 8 x 8
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # current state size. ngf x 16 x 16
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            # current state size. nc x 32 x 32
            # Produce number between -1 and 1, as pixel values have been norma
lized to be between -1 and 1
            nn.Tanh()
        )

    def forward(self, input):
        return self.main(input)

class Discriminator_WGAN(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator_WGAN, self).__init__()
        self.ngpu = ngpu
        ##### YOUR CODE #####
        #####
        self.main = nn.Sequential(
            # input is (nc) x 32 x 32
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 16 x 16
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 8 x 8
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 4 x 4
            nn.Conv2d(ndf * 4, 1, 4, 1, 0, bias=False),
            # state size. (ndf*4) x 1 x 1
            #nn.Sigmoid() # Produce probability
        )
        ##### END YOUR CODE #####
        #####

```

```

def forward(self, input):
    return self.main(input)

netG = initialize_net(Generator, weights_init, device, ngpu)
netD = initialize_net(Discriminator_WGAN, weights_init, device, ngpu)

Generator(
    (main): Sequential(
      (0): ConvTranspose2d(100, 32, kernel_size=(4, 4), stride=(1, 1), bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU(inplace=True)
      (6): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (7): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU(inplace=True)
      (9): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (10): Tanh()
    )
)

Discriminator_WGAN(
    (main): Sequential(
      (0): Conv2d(1, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (1): LeakyReLU(negative_slope=0.2, inplace=True)
      (2): Conv2d(8, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (4): LeakyReLU(negative_slope=0.2, inplace=True)
      (5): Conv2d(16, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (7): LeakyReLU(negative_slope=0.2, inplace=True)
      (8): Conv2d(32, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    )
)

```

```

In [10]: ##### YOUR CODE #####
# Setup RMSprop optimizers for both netG and netD with given learning rate as
`lr_rms`
optimizerD = optim.RMSprop(netD.parameters(), lr=lr_rms)
optimizerG = optim.RMSprop(netG.parameters(), lr=lr_rms)
##### # END YOUR CODE #####
real_label = 1.
fake_label = 0.
# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
n_critic = 5
c = 0.01
dataloader_iter = iter(dataloader)

print("Starting Training Loop...")
num_iters = 1000

for iters in range(num_iters):

    #####
    #
    # (1) Train Discriminator more: minimize -(mean(D(real))-mean(D(fake)))
    #####
    #

    for p in netD.parameters():
        p.requires_grad = True

    for idx_critic in range(n_critic):

        netD.zero_grad()

        try:
            data = next(dataloader_iter)
        except StopIteration:
            dataloader_iter = iter(dataloader)
            data = next(dataloader_iter)

        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        D_real = netD(real_cpu).view(-1)

        noise = torch.randn(b_size, nz, 1, 1, device=device)
        fake = netG(noise)
        D_fake = netD(fake).view(-1)

        ##### YOUR CODE #####
        # Define your loss function for variable `D_loss`
        # Label = torch.full((b_size,), real_label, dtype=torch.float, device=d
evice)
        #D_real = criterion(D_real, Label)
        #D_real.backward()

```

```

#D_fake = netD(fake.detach()).view(-1)
#label.fill_(fake_label)
#D_fake = criterion(D_fake, label)
#D_fake.backward()

D_loss = -(D_real.mean() - D_fake.mean())

# Backpropagate the loss function and upate the optimizer
D_loss.backward()

optimizerD.step()
# Clip the gradient with limit `c` by using `clamp_()` function
for p in netD.parameters():
    p.data.clamp_(-c, c)

##### # END YOUR CODE #####

#####
#
# (2) Update G network: minimize -mean(D(fake)) (Update only once in 5 epochs)
#####
#
for p in netD.parameters():
    p.requires_grad = False

netG.zero_grad()

noise = torch.randn(b_size, nz, 1, 1, device=device)
fake = netG(noise)
D_fake = netD(fake).view(-1)

##### YOUR CODE #####
#
# Define your loss function for variable `G_loss`
G_loss = -(D_fake.mean())

# Backpropagate the loss function and upate the optimizer
G_loss.backward()
optimizerG.step()

##### END YOUR CODE #####

# Output training stats
if iters % 10 == 0:
    print('[%4d/%4d]   Loss_D: %6.4f   Loss_G: %6.4f'
          % (iters, num_iters, D_loss.item(), G_loss.item()))

# Save Losses for plotting later
G_losses.append(G_loss.item())
D_losses.append(D_loss.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 100 == 0):
    with torch.no_grad():

```

```
fake = netG(fixed_noise).detach().cpu()  
img_list.append(vutils.make_grid(fake, padding=2, normalize=True))
```



Starting Training Loop...

[0/1000]	Loss_D: -0.0143	Loss_G: 0.0015
[10/1000]	Loss_D: -0.0412	Loss_G: 0.0160
[20/1000]	Loss_D: -0.0248	Loss_G: 0.0061
[30/1000]	Loss_D: -0.0181	Loss_G: 0.0029
[40/1000]	Loss_D: -0.0151	Loss_G: 0.0005
[50/1000]	Loss_D: -0.0139	Loss_G: 0.0088
[60/1000]	Loss_D: -0.0123	Loss_G: 0.0105
[70/1000]	Loss_D: -0.0134	Loss_G: 0.0164
[80/1000]	Loss_D: -0.0100	Loss_G: 0.0138
[90/1000]	Loss_D: -0.0114	Loss_G: -0.0056
[100/1000]	Loss_D: -0.0105	Loss_G: 0.0083
[110/1000]	Loss_D: -0.0093	Loss_G: 0.0151
[120/1000]	Loss_D: -0.0085	Loss_G: 0.0126
[130/1000]	Loss_D: -0.0093	Loss_G: 0.0053
[140/1000]	Loss_D: -0.0081	Loss_G: -0.0021
[150/1000]	Loss_D: -0.0070	Loss_G: 0.0067
[160/1000]	Loss_D: -0.0073	Loss_G: 0.0082
[170/1000]	Loss_D: -0.0078	Loss_G: 0.0044
[180/1000]	Loss_D: -0.0064	Loss_G: 0.0070
[190/1000]	Loss_D: -0.0067	Loss_G: -0.0007
[200/1000]	Loss_D: -0.0054	Loss_G: -0.0037
[210/1000]	Loss_D: -0.0050	Loss_G: 0.0074
[220/1000]	Loss_D: -0.0054	Loss_G: -0.0008
[230/1000]	Loss_D: -0.0059	Loss_G: -0.0035
[240/1000]	Loss_D: -0.0052	Loss_G: -0.0028
[250/1000]	Loss_D: -0.0051	Loss_G: -0.0030
[260/1000]	Loss_D: -0.0048	Loss_G: 0.0017
[270/1000]	Loss_D: -0.0049	Loss_G: -0.0016
[280/1000]	Loss_D: -0.0045	Loss_G: -0.0050
[290/1000]	Loss_D: -0.0036	Loss_G: 0.0090
[300/1000]	Loss_D: -0.0041	Loss_G: -0.0067
[310/1000]	Loss_D: -0.0048	Loss_G: 0.0048
[320/1000]	Loss_D: -0.0042	Loss_G: 0.0007
[330/1000]	Loss_D: -0.0056	Loss_G: 0.0117
[340/1000]	Loss_D: -0.0044	Loss_G: 0.0002
[350/1000]	Loss_D: -0.0041	Loss_G: 0.0091
[360/1000]	Loss_D: -0.0050	Loss_G: 0.0085
[370/1000]	Loss_D: -0.0040	Loss_G: 0.0090
[380/1000]	Loss_D: -0.0044	Loss_G: -0.0046
[390/1000]	Loss_D: -0.0043	Loss_G: 0.0066
[400/1000]	Loss_D: -0.0043	Loss_G: -0.0003
[410/1000]	Loss_D: -0.0045	Loss_G: -0.0030
[420/1000]	Loss_D: -0.0032	Loss_G: 0.0012
[430/1000]	Loss_D: -0.0037	Loss_G: -0.0035
[440/1000]	Loss_D: -0.0039	Loss_G: -0.0078
[450/1000]	Loss_D: -0.0056	Loss_G: 0.0004
[460/1000]	Loss_D: -0.0037	Loss_G: -0.0032
[470/1000]	Loss_D: -0.0041	Loss_G: -0.0002
[480/1000]	Loss_D: -0.0042	Loss_G: 0.0062
[490/1000]	Loss_D: -0.0040	Loss_G: -0.0030
[500/1000]	Loss_D: -0.0040	Loss_G: 0.0039
[510/1000]	Loss_D: -0.0029	Loss_G: 0.0010
[520/1000]	Loss_D: -0.0037	Loss_G: -0.0098
[530/1000]	Loss_D: -0.0031	Loss_G: 0.0062
[540/1000]	Loss_D: -0.0040	Loss_G: 0.0061
[550/1000]	Loss_D: -0.0040	Loss_G: -0.0065

[560/1000]	Loss_D: -0.0037	Loss_G: -0.0031
[570/1000]	Loss_D: -0.0037	Loss_G: 0.0037
[580/1000]	Loss_D: -0.0038	Loss_G: -0.0070
[590/1000]	Loss_D: -0.0033	Loss_G: -0.0042
[600/1000]	Loss_D: -0.0034	Loss_G: -0.0010
[610/1000]	Loss_D: -0.0034	Loss_G: 0.0004
[620/1000]	Loss_D: -0.0029	Loss_G: -0.0046
[630/1000]	Loss_D: -0.0034	Loss_G: 0.0027
[640/1000]	Loss_D: -0.0033	Loss_G: -0.0074
[650/1000]	Loss_D: -0.0033	Loss_G: -0.0061
[660/1000]	Loss_D: -0.0032	Loss_G: -0.0090
[670/1000]	Loss_D: -0.0025	Loss_G: -0.0066
[680/1000]	Loss_D: -0.0029	Loss_G: -0.0023
[690/1000]	Loss_D: -0.0030	Loss_G: -0.0096
[700/1000]	Loss_D: -0.0031	Loss_G: -0.0079
[710/1000]	Loss_D: -0.0027	Loss_G: -0.0059
[720/1000]	Loss_D: -0.0028	Loss_G: -0.0125
[730/1000]	Loss_D: -0.0027	Loss_G: -0.0034
[740/1000]	Loss_D: -0.0026	Loss_G: -0.0027
[750/1000]	Loss_D: -0.0030	Loss_G: -0.0022
[760/1000]	Loss_D: -0.0026	Loss_G: -0.0004
[770/1000]	Loss_D: -0.0027	Loss_G: -0.0013
[780/1000]	Loss_D: -0.0025	Loss_G: -0.0057
[790/1000]	Loss_D: -0.0025	Loss_G: -0.0010
[800/1000]	Loss_D: -0.0033	Loss_G: -0.0005
[810/1000]	Loss_D: -0.0029	Loss_G: -0.0045
[820/1000]	Loss_D: -0.0025	Loss_G: -0.0014
[830/1000]	Loss_D: -0.0024	Loss_G: -0.0007
[840/1000]	Loss_D: -0.0027	Loss_G: -0.0045
[850/1000]	Loss_D: -0.0035	Loss_G: -0.0030
[860/1000]	Loss_D: -0.0022	Loss_G: 0.0008
[870/1000]	Loss_D: -0.0022	Loss_G: 0.0040
[880/1000]	Loss_D: -0.0029	Loss_G: -0.0048
[890/1000]	Loss_D: -0.0019	Loss_G: -0.0132
[900/1000]	Loss_D: -0.0037	Loss_G: -0.0055
[910/1000]	Loss_D: -0.0027	Loss_G: -0.0050
[920/1000]	Loss_D: -0.0022	Loss_G: 0.0006
[930/1000]	Loss_D: -0.0032	Loss_G: -0.0056
[940/1000]	Loss_D: -0.0023	Loss_G: -0.0085
[950/1000]	Loss_D: -0.0024	Loss_G: -0.0107
[960/1000]	Loss_D: -0.0026	Loss_G: -0.0080
[970/1000]	Loss_D: -0.0023	Loss_G: -0.0011
[980/1000]	Loss_D: -0.0023	Loss_G: -0.0013
[990/1000]	Loss_D: -0.0020	Loss_G: -0.0119


```
In [11]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)
```



(Optional and ungraded) Exercise 3: Implement the WGAN with Gradient Penalty

1. Use slide 20 in [Lecture note for WGAN \(https://www.davidinouye.com/course/ece57000-fall-2021/lectures/wasserstein-gan.pdf\)](https://www.davidinouye.com/course/ece57000-fall-2021/lectures/wasserstein-gan.pdf) to implement WGAN-GP algorithm.
 - A. Use the same discriminator and generator as in Exercise 2.
 - B. Use Adam optimizer for WGAN-GP.
 - C. If implemented correctly, we have setup some hyperparameters (different than the original algorithm) that seem to work in this situation.
 - D. For calculating the gradient penalty term, you will need to:
 - a. Create a batch of interpolated samples.
 - b. Pass this interpolated batch through the discriminator.
 - c. Compute the gradient of the discriminator with respect to the samples using `torch.autograd.grad` (<https://pytorch.org/docs/stable/generated/torch.autograd.grad.html>). You will need to set:
 - i. `outputs`
 - ii. `inputs`
 - iii. `grad_outputs`
 - iv. `create_graph=True` and `retain_graph=True` (because we want to backprop through this gradient calculation for the final objective.)
 - v. Hint: Also make sure to understand the return result of this function to extract the gradients as necessary.
 - d. Compute the gradient penalty (Hint: For numerical stability, we found that `grad_norm = torch.sqrt((grad**2).sum(1) + 1e-14)` is a simple way to compute the norm.)
 - e. Use $\lambda = 10$ for the gradient penalty as in the original paper.
2. Train the model with modified networks and visualize the results.

```
In [20]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
t = torch.tensor([1,2], device=device)
```

```

In [ ]: # Setup networks for WGAN-GP
netG = initialize_net(Generator, weights_init, device, ngpu)
netD = initialize_net(Discriminator_WGAN, weights_init, device, ngpu)

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=5e-4, betas=(0.5, 0.9))
optimizerG = optim.Adam(netG.parameters(), lr=5e-4, betas=(0.5, 0.9))

# Training Loop

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
n_critic = 5
dataloader_iter = iter(dataloader)

print("Starting Training Loop...")
num_iters = 1000

for iters in range(num_iters):

    #####
    #
    # (1) Train Discriminator more: minimize -(mean(D(real))-mean(D(fake)))+GP
    #####
    #

    for p in netD.parameters():
        p.requires_grad = True

    for idx_critic in range(n_critic):

        netD.zero_grad()

        try:
            data = next(dataloader_iter)
        except StopIteration:
            dataloader_iter = iter(dataloader)

        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        D_real = netD(real_cpu).view(-1)

        noise = torch.randn(b_size, nz, 1, 1, device=device)
        fake = netG(noise)
        D_fake = netD(fake).view(-1)

        ##### YOUR CODE #####
        # Compute the gradient penalty term

        # Define your Loss function for variable `D_Loss`

        # Backpropagate the Loss function and upate the optimizer

```

```

##### # END YOUR CODE #####

#####

#
# (2) Update G network: minimize -mean(D(fake)) (Update only once in 5 epochs)
#####
#
for p in netD.parameters():
    p.requires_grad = False

netG.zero_grad()

noise = torch.randn(b_size, nz, 1, 1, device=device)
fake = netG(noise)
D_fake = netD(fake).view(-1)

##### YOUR CODE #####
#
# Define your loss function for variable `G_loss`

# Backpropagate the loss function and update the optimizer

##### END YOUR CODE #####

# Output training stats
if iters % 10 == 0:
    print('[%4d/%4d]   Loss_D: %6.4f   Loss_G: %6.4f'
          % (iters, num_iters, D_loss.item(), G_loss.item()))

# Save Losses for plotting later
G_losses.append(G_loss.item())
D_losses.append(D_loss.item())

# Check how the generator is doing by saving G's output on fixed_noise
if (iters % 100 == 0):
    with torch.no_grad():
        fake = netG(fixed_noise).detach().cpu()
        img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

```

```

In [ ]: # plot the loss for generator and discriminator
plot_GAN_loss([G_losses, D_losses], ["G", "D"])

# Grab a batch of real images from the dataloader
plot_real_fake_images(next(iter(dataloader)), img_list)

```